

oreblocks: License-Free Synthetic Ore-Body Block Models with a Stamped Exact Ultimate-Pit Optimum

Felipe Santibañez-Leal 

Technical note · v1.1, 18 September 2026 · Licence CC BY 4.0

CAOS open-research programme · CAOS_OreBlocks, the oreblocks synthetic block-model package

Version DOI 10.5281/zenodo.22822558 · concept DOI (always latest) 10.5281/zenodo.21512088

Not peer reviewed. Source code, data and computational records: github.com/fsantibanezleal/CAOS_OreBlocks

Abstract. Benchmarking an open-pit optimiser, a mine-dispatch simulator, or a teaching tool needs realistic ore-body block models with a *known* answer, but the standard public instances (the MineLib library) are granted for academic download only and cannot be redistributed, and real mine models are proprietary. We present *oreblocks*, a small pure-Python package that generates synthetic three-dimensional ore-body block models of the *same nature* as the public instances (seeded deposit archetypes with per-block grades, bench structure, slope precedence and net-value economics) and stamps each one with its *exact* ultimate-pit optimum. The optimum is computed by the classical reduction of the ultimate-pit problem to a maximum closure and thence to a minimum s - t cut (Picard’s theorem), solved by a Dinic max-flow with self-verifying closure and value-identity checks, and every instance is emitted in exact MineLib `.blocks/.prec/.upit` format with a sidecar carrying the stamped optimum, so any solver can be validated against a known-by-construction answer without touching licensed data. We validate the solver *independently*: on all four deposit archetypes its pit value agrees with a totally-unimodular linear program solved by a separate method (SciPy HiGHS) to relative 10^{-15} , i.e. to machine precision, and the exact solve runs in under a second up to 32,000 blocks. The package is deterministic, MIT-licensed, and consumed by downstream pit, dispatch and haulage tools; every reported number regenerates from the archived validation run.

Index Terms. Open-pit mining, ultimate pit limit, maximum closure, minimum cut, block model, MineLib, synthetic benchmark, reproducible software, operations research.

I. INTRODUCTION

THE first question of open-pit mine planning is the *ultimate pit*: given a three-dimensional block model in which every block carries a net economic value (its recoverable metal minus the cost to mine and process it) and a set of slope-precedence constraints (a block can be mined only after the blocks above it in the pit wall), which set of blocks maximises total value while respecting the slopes? The problem has a classical exact solution: Lerchs and Grossmann [1] gave a graph algorithm in 1965, and Picard [2] showed that the ultimate pit is the maximum-value closure of the precedence graph, which is the source side of a minimum s - t cut and so

is solvable exactly by any max-flow algorithm [3], [4], [5], [6]. This is textbook operations research [8], and any new pit optimiser, block-scheduling heuristic, or dispatch simulator wants to be tested against it.

The obstacle is data. The community’s shared benchmark, MineLib [7], collects real and realistic instances with published optima, but its instances are granted for academic download only and may not be redistributed; real mine block models are proprietary. A developer who wants to distribute a reproducible test suite, a teaching notebook a student can run, or a continuous-integration job that certifies a solver against a known answer, cannot bundle MineLib data and cannot use a real mine. What is missing is a way to *generate* instances of the same nature, with a known optimum, under a permissive license.

oreblocks fills that gap. It is a small pure-Python package that (i) generates seeded synthetic ore-body block models from four geological archetypes, with per-block grades, bench structure, slope precedence and net-value economics; (ii) solves each instance’s ultimate pit *exactly* by the maximum-closure/min-cut reduction; and (iii) writes the instance in exact MineLib format with a sidecar carrying the stamped optimum, so any downstream solver can be validated against a known-by-construction answer. Everything is deterministic given a seed and every instance is clearly labelled synthetic. A synthetic benchmark is useful only if its stamped answer is correct, so we validate the solver *independently* (Section V) against a linear program solved by an entirely different method, finding agreement to machine precision.

II. STATEMENT OF NEED

The users are developers and educators who need ore-body instances with ground truth and a license that permits redistribution. A pit-optimisation library needs regression tests with known optima; a mine-dispatch or haulage simulator (the downstream consumers of this package) needs geology-grounded scenarios rather than random numbers; a course needs instances a student can generate and solve on a laptop. MineLib serves research but its redistribution terms rule out all three, and real models are unavailable. *oreblocks* provides synthetic instances of the same structure (three-dimensional

F. Santibañez-Leal is with the CAOS open-research programme, Santiago, Chile (e-mail: fsantibanez@gmail.com; ORCID 0000-0002-0150-3246).

Package *oreblocks*; source code under the MIT license at https://github.com/fsantibanezleal/CAOS_OreBlocks.

benches, spatially-correlated grades, slope precedence, net values) whose optimum is not looked up but *computed exactly and stamped*, under the MIT license. The value is a benchmark generator: not a single dataset but a seeded family, each member carrying its own certified answer.

III. SYNTHETIC ORE BODIES

An instance is a block grid (bench levels increasing upward, as in the published MineLib convention) carrying, per block, a grade, a tonnage and a density. Grades come from a deterministic geological trend plus a spatially-correlated (box-smoothed) random field, seeded so an instance is byte-identical given its (archetype, dims, seed). Four archetypes span the common deposit shapes: `porphyry` (a buried ellipsoidal high-grade shell that yields broad bowl pits), `vein` (a dipping tabular zone giving narrow steep pits), `layered` (horizontal stratabound bands), and `core_halo` (a rich core inside a low-grade halo). Each block's net value follows a standard mine-vs-process economics: revenue from recovered metal at a price, less mining and processing costs, with the optimal destination (mill or waste) chosen per block. Slope precedence is generated as a discrete cone (a block requires the pattern of blocks above it within a slope angle), stored in the compressed-sparse-row layout MineLib's `.prec` format uses. The result is a block model of the same nature as a real one, with the geology, the economics and the slope constraints that make the ultimate-pit problem non-trivial.

IV. THE EXACT ULTIMATE-PIT SOLVER

The ultimate pit is the maximum-weight closed set of the precedence graph: a set P of blocks such that every block's predecessors are also in P (the pit is slope-feasible) and $\sum_{b \in P} v_b$ is maximal. Picard's theorem [2] reduces this to a minimum cut (Fig. 1). Build a network on the blocks plus a source S and sink T : an arc $S \rightarrow b$ of capacity v_b for every block with $v_b > 0$, an arc $b \rightarrow T$ of capacity $-v_b$ for every block with $v_b < 0$, and an arc $b \rightarrow p$ of infinite capacity for every precedence relation “ b requires p ”. After a maximum S – T flow, the blocks reachable from S in the residual graph form the optimal pit, and

$$\text{pit value} = \sum_{b: v_b > 0} v_b - \text{maxflow}. \quad (1)$$

The infinite-capacity precedence arcs guarantee that no finite cut ever separates a block from a predecessor, so the source side is always a valid closure. `oreblocks` computes the flow with Dinic's algorithm [4] (an iterative blocking-flow implementation, so deep pits do not overflow the recursion stack), on exact floating-point capacities with an epsilon-guarded residual. The solve is self-verifying: it asserts closure feasibility (every predecessor of an in-pit block is in the pit) and the value identity above before returning, so a silently wrong answer cannot be stamped. This is exactly the max-closure route used by the established exact solvers [5], [6]; the package targets the 10^3 – 10^5 -block instances a synthetic benchmark needs, where a plain exact max-flow is more than fast enough.

TABLE I
STAMPED EXACT OPTIMA OF THE FOUR ARCHETYPES
($24 \times 24 \times 12 = 6,912$ BLOCKS, SEED 11). THE VALUE-IDENTITY GAP IS
THE SOLVER'S OWN SELF-CHECK RESIDUAL. `LAYERED` IS ECONOMIC
THROUGHOUT AT THIS CONFIGURATION, SO ITS PIT IS THE WHOLE
MODEL.

Archetype	Pit value	In pit	Fraction	Gap
porphyry	4.71×10^8	4,631	0.67	6×10^{-8}
vein	3.70×10^8	4,451	0.64	6×10^{-8}
layered	1.46×10^9	6,912	1.00	2×10^{-7}
core_halo	9.43×10^7	3,780	0.55	0

V. VALIDATION

A synthetic benchmark is only as reliable as its stamped answer, so we check the solver against an independent method. The ultimate-pit problem also has a linear-programming formulation: maximise $\sum_b v_b x_b$ subject to $x_b \leq x_p$ for every precedence relation and $0 \leq x_b \leq 1$. Its constraint matrix is a network (incidence) matrix and hence totally unimodular, so the LP relaxation is integral and its optimum equals the exact ultimate-pit value. We solve this LP with SciPy's HiGHS backend, a completely different algorithm from the max-flow solver, on every archetype and compare (Fig. 2a). The two agree to relative 10^{-15} (machine precision) on all four archetypes: 5.7×10^{-16} (porphyry), 4.9×10^{-16} (vein), 1.7×10^{-15} (layered), 1.4×10^{-16} (core_halo). This is an end-to-end check of the whole solve (the network construction, the flow, the residual reachability, and the value identity), by a method that shares no code with it. It corroborates the package's independent cross-check against a separate min-cut engine that reproduces the published MineLib optima (newman1, zuck_small, kd).

The exact solve is fast (Fig. 2b): from 0.02 s at 2,048 blocks to 0.92 s at 32,000 blocks, comfortably within the range a synthetic benchmark occupies. Table I lists the stamped optima of the four archetypes at a common size, each with the solver's own value-identity gap (the residual of the identity above), which is at most 2.4×10^{-7} , a relative residual of order 10^{-16} of the pit value (double-precision round-off), confirming the internal consistency of every stamped answer. Figure 3 shows what an instance looks like: a porphyry cross-section with its exact pit, a slope-feasible bowl that steps down bench by bench to the high-grade core.

VI. MINELIB FORMAT AND THE STAMPED TWIN

An instance is emitted as a MineLib triplet plus a sidecar. The `.blocks` file lists per-block coordinates and attributes, `.prec` the precedence in the compressed-sparse-row form MineLib uses, and `.upit` the per-block economic value, so the triplet is byte-consumable by any MineLib reader. The `.meta.json` sidecar carries the economics, the slope angle, and the *stamped optimum* (pit value and block count), so a solver under test reads the instance, computes its own pit, and compares against the stamped answer with no licensed data anywhere in the loop. A whole instance with its certified optimum is one call:




Fig. 1. The exact ultimate-pit solve. **Left:** a block model with per-block net values and slope precedence (a block is mineable only after its slope predecessors above it). **Right:** Picard's reduction [2]. Add a source S with an arc to every positive block (capacity v_b), an arc from every negative block to a sink T (capacity $-v_b$), and an infinite-capacity arc for every precedence relation. The optimal pit is the source side of the minimum $s-t$ cut, and its value is the sum of positive block values minus the maximum flow. The infinite-capacity precedence arcs make any finite cut respect the slopes exactly.




Fig. 2. Solver validation, from the archived validation run. **(a)** Independent cross-check: on every archetype the max-closure pit value agrees with the same instance solved as a linear program by SciPy HiGHS (an entirely different algorithm) to relative $\sim 10^{-15}$, machine precision (dashed line). **(b)** The exact solve scales sub-second: 0.02 s at 2,048 blocks to 0.92 s at 32,000 blocks.

```
from oreblocks import make_twin
tw = make_twin("porphyry", dims=(20,20,10),
               seed=42)
# the exact optimum, stamped
tw.upit.pit_value, tw.upit.n_in_pit
# -> .blocks / .prec / .upit / .meta.json
tw.write("out/")
```

VII. ARCHITECTURE AND AVAILABILITY

The package consists of small pure-Python modules (grid, fields, economics, precedence, upit, minelib_io, twins) over NumPy storage, with no compiled dependency, so it installs and runs anywhere Python does. It is deterministic given a seed, unit-tested (the

generators, the precedence, the solver's closure and identity checks, the MineLib round-trip), and MIT-licensed. The downstream CAOS tools (a haulage simulator, a pit-forge twin generator, a dispatch bench) consume it for geology-grounded, license-free instances. Source code and the reproducible validation are available at https://github.com/fsantibanezleal/CAOS_OreBlocks.

VIII. LIMITATIONS

oreblocks generates instances of the *ultimate-pit* problem (the single-period pit limit), not the constrained production-scheduling problem (CPIT/PCPSP) with capacities and time periods [7]; those are a different, harder class.




Fig. 3. A vertical cross-section of a porphyry twin: the grade field (the buried high-grade shell) with the exact ultimate-pit boundary overlaid. The pit is a slope-feasible bowl whose walls step down bench by bench (the discrete slope precedence), bottoming under the high-grade core, exactly the structure a real pit optimiser must reproduce.

The deposits are geologically plausible synthetic fields, not calibrated to any real ore body, and are labelled synthetic throughout. The economics is a standard single-metal mine-vs-process model; multi-element or blending economics would need extension. The exact solver targets the 10^3 – 10^5 -block range a benchmark generator occupies; for the multi-million-block industrial models the specialised pseudoflow implementations [6] are the tool of choice, and `oreblocks` writes MineLib format precisely so such solvers can consume its instances. The package labels every instance synthetic and states what its stamped optimum does and does not certify.

IX. CONCLUSION

`oreblocks` turns the shortage of redistributable ore-body benchmarks into a generator: seeded, synthetic, geologically-structured block models of the MineLib nature, each stamped with its exact ultimate-pit optimum by the classical max-closure/min-cut solve, emitted in MineLib format with the answer in a sidecar. The stamped answer is checked against an independent linear program, with agreement to machine precision. The value is a license-free, reproducible source of realistic pit-optimisation instances with certified optima, for testing solvers, grounding simulators, and teaching, without touching licensed data.

APPENDIX A REPRODUCTION

The validation is `manuscripts/ore-body-twins/figures/upit_validation.py`: it builds each archetype, solves the ultimate pit by the package’s

max-closure solver and independently as a HiGHS linear program, records the agreement, the four stamped optima and the solve-time scaling, and extracts a porphyry cross-section. The figure script reads those artifacts; the max-closure schematic is a hand-authored theme-aware drawing. Re-running reproduces the cross-check (10^{-15} agreement), Table I, and Figs. 2–3; the package is seeded and deterministic.

REFERENCES

- [1] H. Lerchs and I. F. Grossmann, “Optimum design of open-pit mines,” *Transactions of the Canadian Institute of Mining and Metallurgy*, vol. 68, pp. 17–24, 1965.
- [2] J.-C. Picard, “Maximal closure of a graph and applications to combinatorial problems,” *Management Science*, vol. 22, no. 11, pp. 1268–1272, 1976. doi:10.1287/mnsc.22.11.1268.
- [3] L. R. Ford and D. R. Fulkerson, “Maximal flow through a network,” *Canadian Journal of Mathematics*, vol. 8, pp. 399–404, 1956. doi:10.4153/CJM-1956-045-5.
- [4] Y. Diniz, “Dinitz’ algorithm: The original version and Even’s version,” in *Theoretical Computer Science: Essays in Memory of Shimon Even*, LNCS 3895. Berlin: Springer, 2006, pp. 218–240. doi:10.1007/11685654_10.
- [5] D. S. Hochbaum and A. Chen, “Performance analysis and best implementations of old and new algorithms for the open-pit mining problem,” *Operations Research*, vol. 48, no. 6, pp. 894–914, 2000. doi:10.1287/opre.48.6.894.12392.
- [6] D. S. Hochbaum, “The pseudoflow algorithm: A new algorithm for the maximum-flow problem,” *Operations Research*, vol. 56, no. 4, pp. 992–1009, 2008. doi:10.1287/opre.1080.0524.
- [7] D. Espinoza, M. Goycoolea, E. Moreno, and A. Newman, “MineLib: A library of open pit mining problems,” *Annals of Operations Research*, vol. 206, no. 1, pp. 93–114, 2013. doi:10.1007/s10479-012-1258-3.
- [8] A. M. Newman, E. Rubio, R. Caro, A. Weintraub, and K. Eurek, “A review of operations research in mine planning,” *Interfaces*, vol. 40, no. 3, pp. 222–245, 2010. doi:10.1287/inte.1090.0492.